



# 14-848 Cloud Infrastructure

---

APACHE KAFKA

# Agenda

---

- Announcements
- User-response Architectures
- Apache Kafka
- Kafka & Microservices
- Confluent-Kafka




# Announcements

---

- Course Project is Due by this Thursday.
  - Plan to submit ahead of time to avoid last-minute issues
  - Read the submission guidelines carefully
- HW-7 is released on Wednesday.
  - HW-7 is due by November 21<sup>st</sup>

# Request-Response Architecture

- Your Machine Learning Infrastructure offers the ability to train models based on incoming training requests.
- What happens if the number of training requests increases?
  - Your Machine Learning infrastructure can leverage distributed training environment to process multiple requests simultaneously.
- Well, what happens if the number of user requests continue to increase beyond your ML infrastructure resources?






# Request-Response Architecture

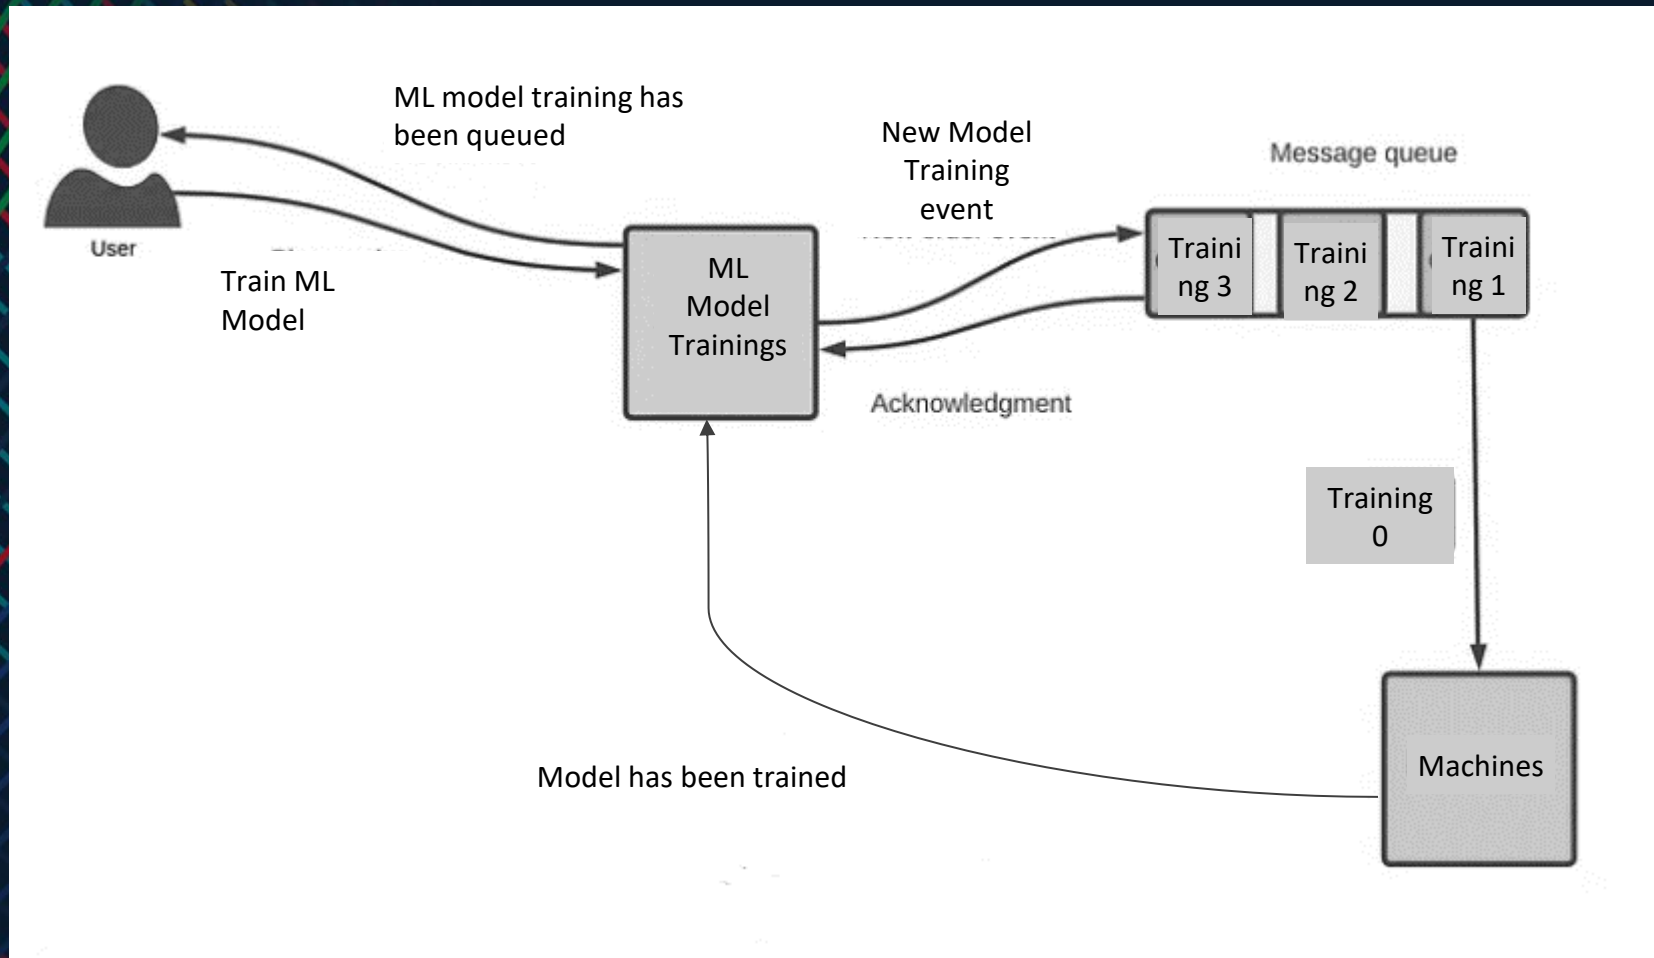
- Well, what happens if the number of user requests continue to increase beyond your Kubernetes cluster resource capability?

**If the number of requests is beyond the available hardware resources, user requests will be dropped because the ML infrastructure can't handle all the incoming requests**

**In other words, your ML infrastructure won't have enough memory or processing power to conduct distribute training**



# Solution: Use Message Queues



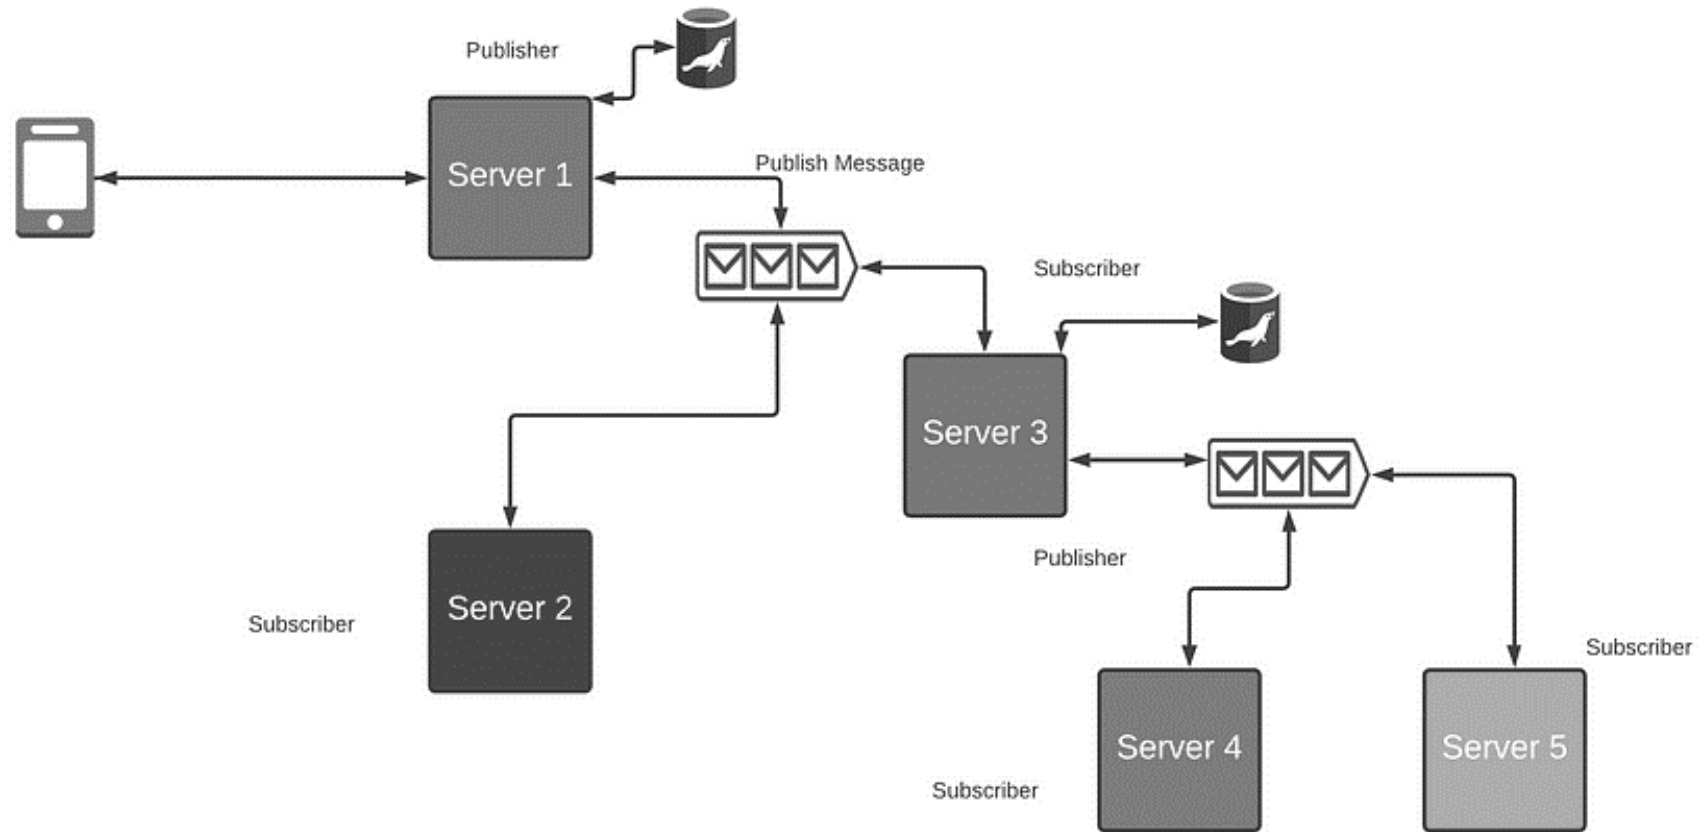




# Apache Kafka

Apache Kafka is a scalable message queue with publish-subscribe software support that offers distributed streaming processing.

# Digression 1: What is a Publish-Subscribe Model?





# Kafka History

---

- Apache Kafka was originally developed by LinkedIn in 2010.
- Later, it was donated to the Apache Software Foundation and was made public in 2011.
- Currently, it is maintained by Confluent under Apache Software Foundation.





# Kafka Advantages

---

- Very high performance: Kafka uses in-memory writes and reads
- High-throughput: LinkedIn leverages Kafka for up-to 3.2M messages/second
- Elastically scalable
- Low operational overhead
- Durable, highly available
- Fault-tolerant: Kafka is highly available and resilient to node failures and supports automatic recovery



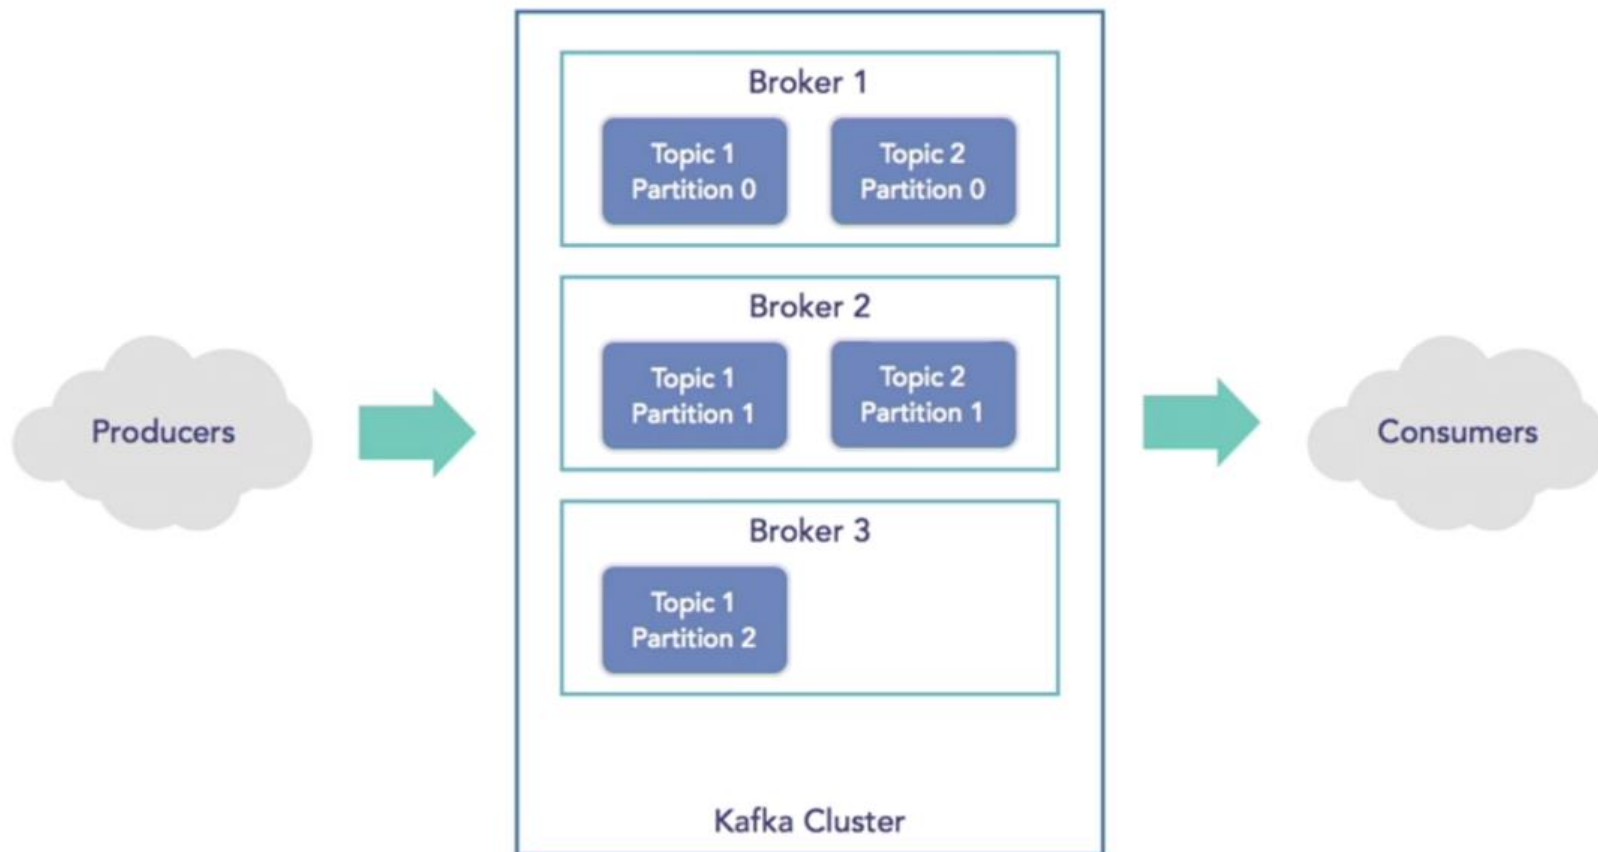
# Kafka Use Cases

---

- **LinkedIn** uses Kafka to prevent spam, collect user interactions to make better connection recommendations in real-time. Kafka is used to process 220B msg/day with a peak of **3.2M msg/second**.
- **Netflix** uses Kafka to apply recommendations in real-time while you are watching the TV shows.
- **X** uses Kafka as part of their Storm real-time data pipelines.
- **Spotify** leverages Kafka for log delivery.
- **Loggly** retains Kafka for log collection and processing
- **Other companies:** Airbnb, Cisco, Gnip, InfoChimps, Ooyala, Square, and Uber



# Kafka Architecture





# Topics

---

- Topic is a particular stream of data
- You may have as many topics as you want
- A topic is identified by its name
- A topic is divided into partitions

A decorative plaid pattern in the top-left corner, featuring a grid of red, green, and blue lines on a dark blue background.

# Partitions

---

- Each partition is ordered
- Each message within a partition gets an incremental id, called offset





# Brokers

---

- A Kafka cluster is composed of multiple brokers (servers)
- Each broker is identified with its ID (integer)
- Each broker contains certain topic partitions
- After connecting to any broker, you will be connected to the entire cluster



# Producers

---

- Producers write data to topics (which are made of partitions)
- Producers automatically know to which broker and partition they should write to
- In case of Broker failure, producers will automatically recover the failed broker.
- Producers choose whether to receive **acknowledgement** of data writes (ack)
- Producers choose whether to send a **key** with the message



# Consumers

---

- Consumers read data from a topic (identified by a topic name)
- Consumers know which broker to read from.
- In case of Broker failures, consumers know how to recover the failed broker.
- Data are read in order within each partition.





# Zookeeper

---

- Zookeeper manages brokers (keeps a list of them)
- Zookeeper helps in performing election for partitions
- Zookeeper sends notifications to Kafka in case of changes (e.g., new topic, broker failure, etc.).
- Zookeeper is an important component for Kafka.



# Topic Partitions



# Topic Partitions

---

- A topic partition is the **unit of parallelism** in Apache Kafka.
- **Producers and Brokers:**
  - Writes to different partitions can be done in parallel.
  - Parallelism frees up hardware resources for operations like compression.
- **Consumers:**
  - Consumers can be grouped together into **several consumer groups**.
  - You can have up to one consumer instance per partition within a consumer group.
  - Any additional consumers beyond the number of partitions will remain idle.
- **Kafka Cluster:** More partitions in a Kafka cluster lead to higher system throughput.





## More on Topic Partitions

---

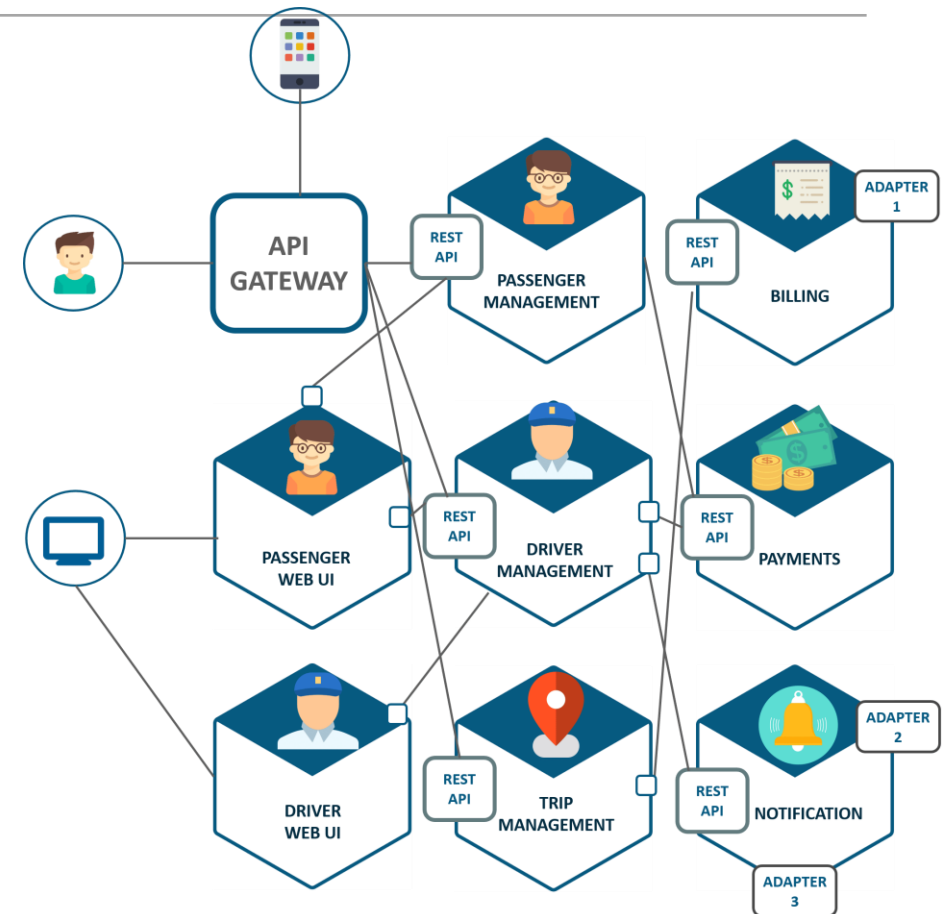
- **Scalable Partitioning:** You can expand the number of partitions as your data needs grow.
- **Keyed Messages Caution:** Exercise caution when messages are produced with keys, as they are deterministically assigned to partitions.
- **Consistent Routing:** Kafka ensures that messages with identical keys are consistently routed to the same partition, preserving the integrity of order-dependent applications.



But ..  
How does Kafka fit in the  
development of  
Large-scale applications?

## Digression 2: Microservices

- Microservices are common software design pattern for applications running at scale.
- Divide a computation into small, mostly stateless components that can be:
  - Easily replicated for scale
  - Communicate with simple protocols
  - Computation is as a swarm of communicating workers.






# Microservices in the Cloud

Typically run as containers using a service deployment and management service on systems like:

- Amazon Elastic Container Service
- Google Kubernetes
- DCOS from Berkeley/Mesosphere
- Docker Swarm





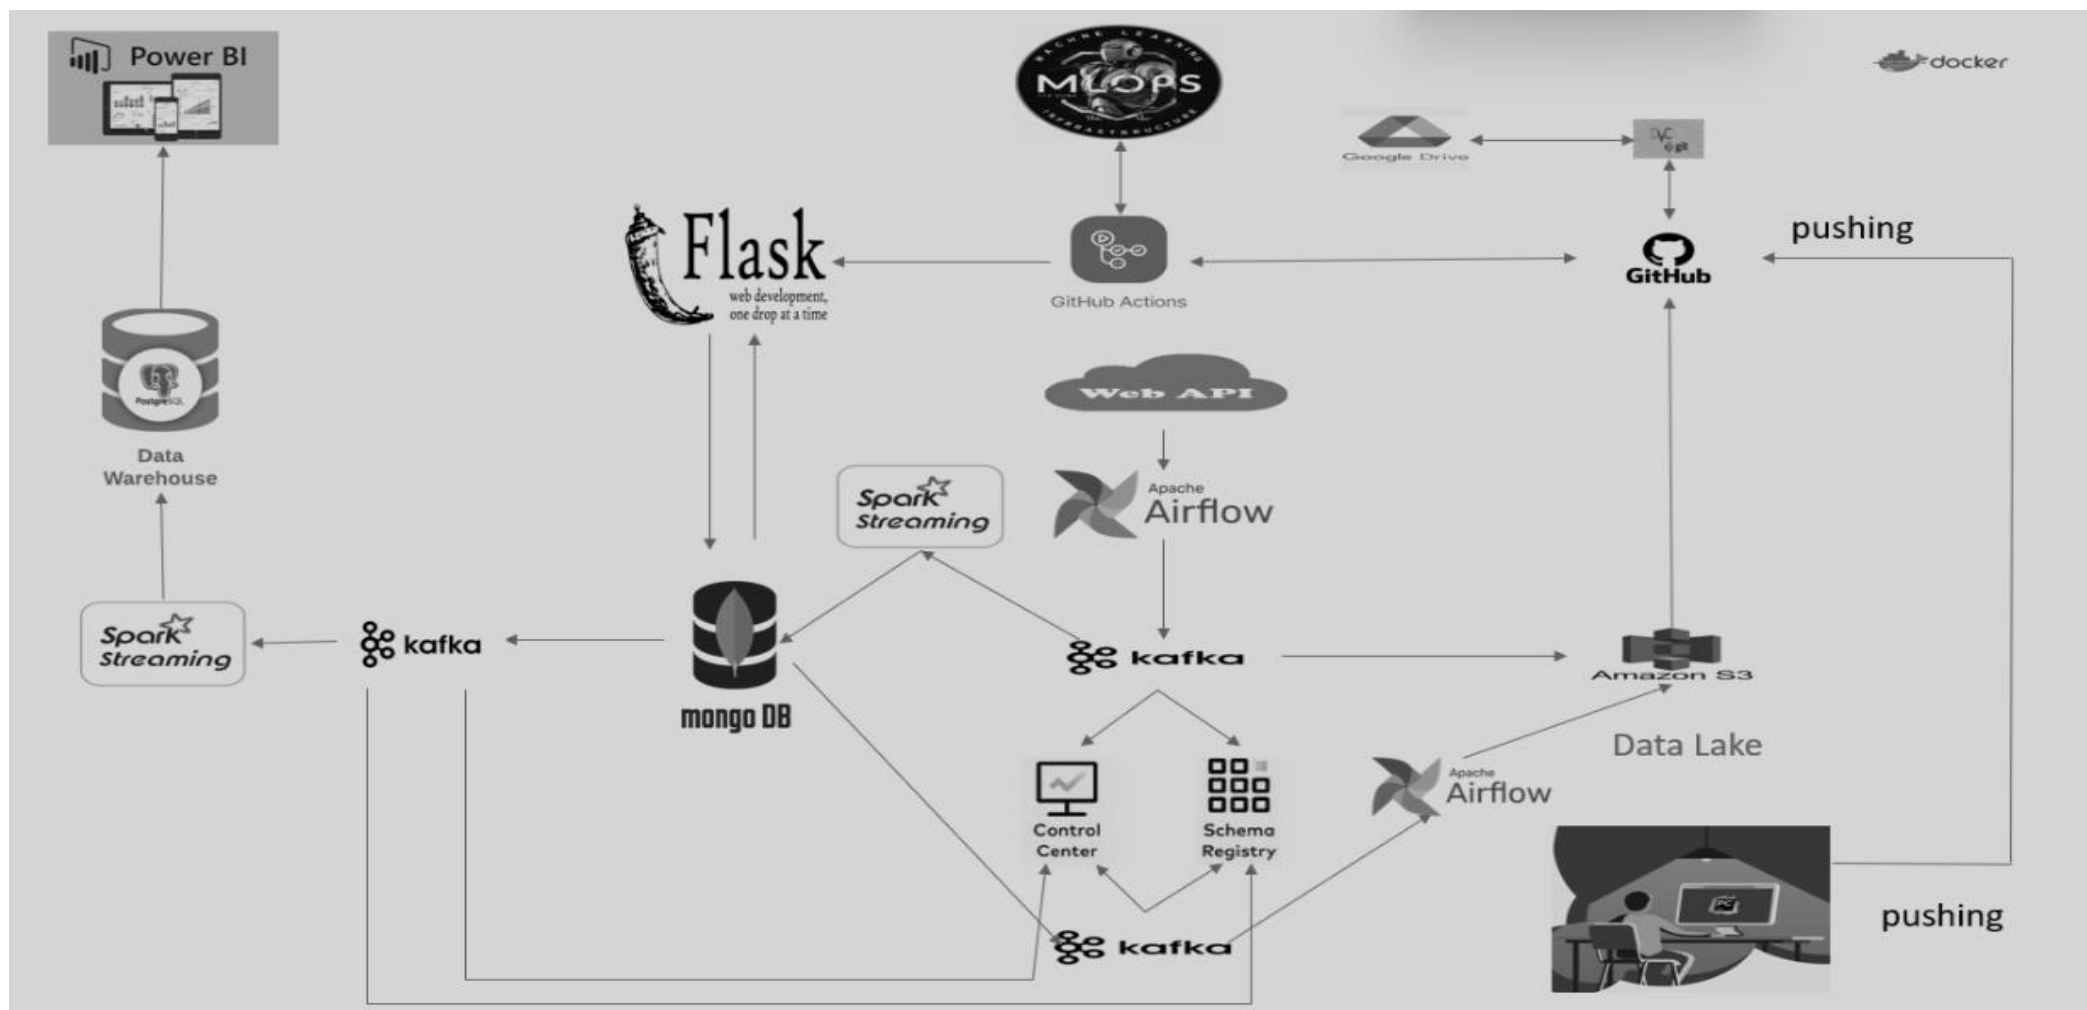


## Digression 3 : REST APIs

---

- REST APIs provide an interface for distributed hypermedia systems, created by Roy Fielding in 2000
- The key abstraction of information in REST is a resource
  - Identified by a resource identifier, ie URI
- Resources can be retrieved or transformed to another state by a set of methods
  - GET/PUT/POST/DELETE/PATCH
- The clients and servers exchange representations of resources by using a standardized interface and protocol typically HTTP

# Kafka and Microservices Combined Example





# Apache Kafka & Confluent Kafka





# Apache Kafka and Confluent Kafka

---

- Handling Kafka requires managing low-level, complex data infrastructure.
- Confluent Kafka is built on top of Apache Kafka to offer complete, fully managed, cloud-native data streaming that's available wherever your data and applications reside.



# Apache Kafka vs Confluent Kafka

	Apache Kafka	Confluent
<b>Serverless</b> Automated, fully managed Kafka clusters with zero ops	—	✓
<b>Elastic Scaling</b> Scale up and down from 0 to GBps without over-provisioning infra	—	✓
<b>Infinite Storage / Tiered Storage</b> Cost-effectively retain data at any scale without growing compute	—	✓
<b>High Availability</b> Guaranteed 99.99% uptime SLA with built-in failover and multi-AZ replication	—	✓
<b>No ZooKeeper management</b> Metadata management completely abstracted away	—	✓
<b>No-touch patching and upgrades</b> Fully optimized infra with zero-downtime patching and upgrades	—	✓

# Take-Home Activity (Complete before next lecture)

---

## Create a “Self-managed” Confluent-Kafka on Your Machine

- Follow the following tutorials:
  - <https://docs.confluent.io/platform/current/get-started/platform-quickstart.html>
- Optional: Deploy Confluent Kafka to GKE:
  - <https://docs.confluent.io/operator/current/co-quickstart.html>